

```
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.datasets import imdb
(train_data, train_labels), (test_data, test_labels) = imdb.load_data(
    num_words=10000)
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/imdb.npz>  
 17464789/17464789 ————— 2s 0us/step

The above code imports tensorflow, keras, and the imdb dataset from keras. It also nick names tensorflow "tf".

```
max([max(sequence) for sequence in train_data])
9999
```

This shows the maximum size of each piece of data

```
import numpy as np
```

This imports numpy as np

```
import numpy as np
def vectorize_sequences(sequences, dimension=10000):
    results = np.zeros((len(sequences), dimension))
    for i, sequence in enumerate(sequences):
        for j in sequence:
            results[i, j] = 1.
    return results
x_train = vectorize_sequences(train_data)
x_test = vectorize_sequences(test_data)

y_train = np.asarray(train_labels).astype("float32")
y_test = np.asarray(test_labels).astype("float32")
```

This encodes the integer sequences by using hot encoding

```
from tensorflow import keras
from tensorflow.keras import layers, regularizers

model = keras.Sequential([
    layers.Dense(16, activation="relu"),
    layers.Dense(16, activation="relu"),
    layers.Dense(1, activation="sigmoid")
])
```

This tells the computer what kind of model we want to use. Here you decide things like the amount of layers you want and what activation function you would like to use

```
x_val = x_train[:10000]
partial_x_train = x_train[10000:]
y_val = y_train[:10000]
partial_y_train = y_train[10000:]
```

This creates a validation set to help train the model

```
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])
```

This is the code that compiles the model

```
history = model.fit(partial_x_train,
                    partial_y_train,
                    epochs=20,
                    batch_size=512,
                    validation_data=(x_val, y_val))
```

```
Epoch 1/20
30/30 _____ 4s 96ms/step - accuracy: 0.6880 - loss:
0.6101 - val_accuracy: 0.8696 - val_loss: 0.4015
Epoch 2/20
30/30 _____ 1s 38ms/step - accuracy: 0.8895 - loss:
0.3528 - val_accuracy: 0.8598 - val_loss: 0.3492
Epoch 3/20
30/30 _____ 2s 49ms/step - accuracy: 0.9154 - loss:
0.2567 - val_accuracy: 0.8845 - val_loss: 0.2920
Epoch 4/20
30/30 _____ 2s 39ms/step - accuracy: 0.9341 - loss:
0.2013 - val_accuracy: 0.8881 - val_loss: 0.2771
Epoch 5/20
30/30 _____ 2s 50ms/step - accuracy: 0.9479 - loss:
0.1643 - val_accuracy: 0.8870 - val_loss: 0.2776
Epoch 6/20
30/30 _____ 2s 39ms/step - accuracy: 0.9565 - loss:
0.1377 - val_accuracy: 0.8870 - val_loss: 0.2918
Epoch 7/20
30/30 _____ 1s 40ms/step - accuracy: 0.9591 - loss:
0.1259 - val_accuracy: 0.8853 - val_loss: 0.3017
Epoch 8/20
30/30 _____ 2s 62ms/step - accuracy: 0.9722 - loss:
0.0989 - val_accuracy: 0.8665 - val_loss: 0.3612
Epoch 9/20
30/30 _____ 2s 37ms/step - accuracy: 0.9708 - loss:
0.0984 - val_accuracy: 0.8802 - val_loss: 0.3276
```

```

Epoch 10/20
30/30 _____ 1s 36ms/step - accuracy: 0.9786 - loss:
0.0753 - val_accuracy: 0.8757 - val_loss: 0.3749
Epoch 11/20
30/30 _____ 1s 37ms/step - accuracy: 0.9849 - loss:
0.0647 - val_accuracy: 0.8784 - val_loss: 0.3632
Epoch 12/20
30/30 _____ 1s 36ms/step - accuracy: 0.9874 - loss:
0.0533 - val_accuracy: 0.8776 - val_loss: 0.3890
Epoch 13/20
30/30 _____ 1s 37ms/step - accuracy: 0.9889 - loss:
0.0483 - val_accuracy: 0.8775 - val_loss: 0.4056
Epoch 14/20
30/30 _____ 1s 38ms/step - accuracy: 0.9934 - loss:
0.0379 - val_accuracy: 0.8735 - val_loss: 0.4277
Epoch 15/20
30/30 _____ 1s 37ms/step - accuracy: 0.9956 - loss:
0.0326 - val_accuracy: 0.8725 - val_loss: 0.4531
Epoch 16/20
30/30 _____ 1s 37ms/step - accuracy: 0.9948 - loss:
0.0288 - val_accuracy: 0.8729 - val_loss: 0.4809
Epoch 17/20
30/30 _____ 2s 60ms/step - accuracy: 0.9970 - loss:
0.0237 - val_accuracy: 0.8710 - val_loss: 0.4945
Epoch 18/20
30/30 _____ 2s 36ms/step - accuracy: 0.9985 - loss:
0.0167 - val_accuracy: 0.8577 - val_loss: 0.6314
Epoch 19/20
30/30 _____ 1s 37ms/step - accuracy: 0.9932 - loss:
0.0262 - val_accuracy: 0.8475 - val_loss: 0.6525
Epoch 20/20
30/30 _____ 1s 36ms/step - accuracy: 0.9958 - loss:
0.0208 - val_accuracy: 0.8690 - val_loss: 0.5661

```

This trains the model

```


history_dict = history.history
history_dict.keys()

dict_keys(['accuracy', 'loss', 'val_accuracy', 'val_loss'])

import matplotlib.pyplot as plt
history_dict = history.history
loss_values = history_dict["loss"]
val_loss_values = history_dict["val_loss"]
epochs = range(1, len(loss_values) + 1)
plt.plot(epochs, loss_values, "bo", label="Training loss")
plt.plot(epochs, val_loss_values, "b", label="Validation loss")
plt.title("Training and validation loss")


```

```
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.show()
```



This creates a graph showing the training and validation loss

```
plt.clf()
acc = history_dict["accuracy"]
val_acc = history_dict["val_accuracy"]
plt.plot(epochs, acc, "bo", label="Training acc")
plt.plot(epochs, val_acc, "b", label="Validation acc")
plt.title("Training and validation accuracy")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()
plt.show()
```



This creates a graph showing the training and validation accuracy

It should be noted in order to make this file more readable I have not included separate code for each time that I have run the The training and validation loss/accuracy graphs. I used the above code and modified the model it runs off of to get the correct epoch number for each model.

```
model = keras.Sequential([
    layers.Dense(16, activation="relu"),
    layers.Dense(16, activation="relu"),
    layers.Dense(1, activation="sigmoid")
])
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])
model.fit(x_train, y_train, epochs=5, batch_size=512)
results = model.evaluate(x_test, y_test)
```

Epoch 1/5  
 49/49 ————— 3s 27ms/step - accuracy: 0.7048 - loss: 0.6099  
 Epoch 2/5  
 49/49 ————— 3s 40ms/step - accuracy: 0.8890 - loss: 0.3441

```

Epoch 3/5
49/49 _____ 2s 37ms/step - accuracy: 0.9133 - loss:
0.2492
Epoch 4/5
49/49 _____ 2s 27ms/step - accuracy: 0.9338 - loss:
0.1970
Epoch 5/5
49/49 _____ 1s 27ms/step - accuracy: 0.9399 - loss:
0.1734
782/782 _____ 2s 3ms/step - accuracy: 0.8834 - loss:
0.2866

```

With the original model set up from class the accuracy is 0.8852 and the loss is 0.2850

```

model = keras.Sequential([
    layers.Dense(16, activation="relu"),
    layers.Dense(16, activation="relu"),
    layers.Dense(16, activation="relu"),
    layers.Dense(1, activation="sigmoid")
])
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])
model.fit(x_train, y_train, epochs=3, batch_size=512)
results = model.evaluate(x_test, y_test)

Epoch 1/3
49/49 _____ 4s 30ms/step - accuracy: 0.6992 - loss:
0.5763
Epoch 2/3
49/49 _____ 3s 45ms/step - accuracy: 0.8993 - loss:
0.2896
Epoch 3/3
49/49 _____ 3s 45ms/step - accuracy: 0.9214 - loss:
0.2154
782/782 _____ 3s 4ms/step - accuracy: 0.8859 - loss:
0.2823

```

With adding another hidden layer in the model definition step the training and validation accuracy and training and validation loss graphs showed 3 epochs as the best number to use so I changed that value when retraining the model from scratch. And the accuracy is 0.8832 and the loss is 0.2854

```

model = keras.Sequential([
    layers.Dense(64, activation="relu"),
    layers.Dense(64, activation="relu"),
    layers.Dense(1, activation="sigmoid")
])
model.compile(optimizer="rmsprop",

```

```

        loss="binary_crossentropy",
        metrics=["accuracy"])
model.fit(x_train, y_train, epochs=2, batch_size=512)
results = model.evaluate(x_test, y_test)

Epoch 1/2
49/49 _____ 4s 48ms/step - accuracy: 0.7179 - loss:
0.5409
Epoch 2/2
49/49 _____ 2s 47ms/step - accuracy: 0.8999 - loss:
0.2665
782/782 _____ 4s 5ms/step - accuracy: 0.8695 - loss:
0.3152

```

With using layers with more hidden units in the model definition step the training and validation accuracy and training and validation loss graphs showed 2 epochs as the best number to use so I changed that value when retraining the model from scratch. And the accuracy is 0.8782 and the loss is 0.2939

```

model = keras.Sequential([
    layers.Dense(16, activation="relu"),
    layers.Dense(16, activation="relu"),
    layers.Dense(1, activation="sigmoid")
])
model.compile(optimizer="rmsprop",
              loss="mse",
              metrics=["accuracy"])
model.fit(x_train, y_train, epochs=3, batch_size=512)
results = model.evaluate(x_test, y_test)

Epoch 1/3
49/49 _____ 2s 27ms/step - accuracy: 0.7290 - loss:
0.1937
Epoch 2/3
49/49 _____ 3s 31ms/step - accuracy: 0.9003 - loss:
0.0888
Epoch 3/3
49/49 _____ 3s 31ms/step - accuracy: 0.9222 - loss:
0.0677
782/782 _____ 3s 3ms/step - accuracy: 0.8866 - loss:
0.0855

```

With using the "mse" loss function instead of "binary\_crossentropy" in the model definition step the training and validation accuracy and training and validation loss graphs showed 3 epochs as the best number to use so I changed that value when retraining the model from scratch. And the accuracy is 0.8819 and the loss is 0.0894

```

model = keras.Sequential([
    layers.Dense(16, activation="tanh"),

```

```

        layers.Dense(16, activation="tanh"),
        layers.Dense(1, activation="sigmoid")
    ])
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])
model.fit(x_train, y_train, epochs=3, batch_size=512)
results = model.evaluate(x_test, y_test)

Epoch 1/3
49/49 ————— 3s 31ms/step - accuracy: 0.7410 - loss:
0.5322
Epoch 2/3
49/49 ————— 3s 30ms/step - accuracy: 0.9090 - loss:
0.2539
Epoch 3/3
49/49 ————— 1s 27ms/step - accuracy: 0.9296 - loss:
0.1932
782/782 ————— 2s 3ms/step - accuracy: 0.8788 - loss:
0.3021

```

With using the "tanh" activation instead of "relu" in the model definition step the training and validation accuracy and training and validation loss graphs showed 3 epochs as the best number to use so I changed that value when retraining the model from scratch. And the accuracy is 0.8868 and the loss is 0.2857

```

model = keras.Sequential([
    layers.Dense(16, activation="relu",
kernel_regularizer=regularizers.L2(0.01)),
    layers.Dense(16, activation="relu",
kernel_regularizer=regularizers.L2(0.01)),
    layers.Dense(1, activation="sigmoid",
kernel_regularizer=regularizers.L2(0.01))
])
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])
model.fit(x_train, y_train, epochs=16, batch_size=512)
results = model.evaluate(x_test, y_test)

Epoch 1/16
49/49 ————— 3s 38ms/step - accuracy: 0.7307 - loss:
0.9170
Epoch 2/16
49/49 ————— 1s 27ms/step - accuracy: 0.8846 - loss:
0.5708

```



```

Epoch 3/16
49/49 _____ 3s 39ms/step - accuracy: 0.8927 - loss:
0.5114
Epoch 4/16
49/49 _____ 2s 27ms/step - accuracy: 0.8984 - loss:
0.4833
Epoch 5/16
49/49 _____ 3s 28ms/step - accuracy: 0.9001 - loss:
0.4616
Epoch 6/16
49/49 _____ 1s 27ms/step - accuracy: 0.8971 - loss:
0.4540
Epoch 7/16
49/49 _____ 1s 27ms/step - accuracy: 0.9013 - loss:
0.4393
Epoch 8/16
49/49 _____ 3s 27ms/step - accuracy: 0.8972 - loss:
0.4374
Epoch 9/16
49/49 _____ 2s 42ms/step - accuracy: 0.9008 - loss:
0.4301
Epoch 10/16
49/49 _____ 2s 38ms/step - accuracy: 0.8994 - loss:
0.4263
Epoch 11/16
49/49 _____ 2s 26ms/step - accuracy: 0.9016 - loss:
0.4200
Epoch 12/16
49/49 _____ 1s 27ms/step - accuracy: 0.9009 - loss:
0.4199
Epoch 13/16
49/49 _____ 2s 26ms/step - accuracy: 0.9044 - loss:
0.4093
Epoch 14/16
49/49 _____ 1s 27ms/step - accuracy: 0.9101 - loss:
0.4036
Epoch 15/16
49/49 _____ 1s 27ms/step - accuracy: 0.8976 - loss:
0.4162
Epoch 16/16
49/49 _____ 3s 41ms/step - accuracy: 0.9003 - loss:
0.4136
782/782 _____ 2s 3ms/step - accuracy: 0.8826 - loss:
0.4459

```

With using the the L2 regularization in the model definition step the training and validation accuracy and training and validation loss graphs showed 16 epochs as the best number to use so I changed that value when retraining the model from scratch. And the accuracy is 0.8712 and the loss is 0.4673

```
history = model.fit(partial_x_train,
                    partial_y_train,
                    epochs=20,
                    batch_size=512,
                    validation_data=(x_val, y_val))
```

Epoch 1/20

30/30 \_\_\_\_\_ 2s 60ms/step - accuracy: 0.9123 - loss: 0.3951 - val\_accuracy: 0.9053 - val\_loss: 0.4013

Epoch 2/20

30/30 \_\_\_\_\_ 3s 60ms/step - accuracy: 0.9180 - loss: 0.3857 - val\_accuracy: 0.8965 - val\_loss: 0.4188

Epoch 3/20

30/30 \_\_\_\_\_ 1s 38ms/step - accuracy: 0.9219 - loss: 0.3787 - val\_accuracy: 0.8720 - val\_loss: 0.4638

Epoch 4/20

30/30 \_\_\_\_\_ 1s 36ms/step - accuracy: 0.9242 - loss: 0.3755 - val\_accuracy: 0.8846 - val\_loss: 0.4422

Epoch 5/20

30/30 \_\_\_\_\_ 1s 35ms/step - accuracy: 0.9249 - loss: 0.3712 - val\_accuracy: 0.8868 - val\_loss: 0.4363

Epoch 6/20

30/30 \_\_\_\_\_ 1s 34ms/step - accuracy: 0.9267 - loss: 0.3742 - val\_accuracy: 0.8802 - val\_loss: 0.4474

Epoch 7/20

30/30 \_\_\_\_\_ 1s 35ms/step - accuracy: 0.9217 - loss: 0.3761 - val\_accuracy: 0.8835 - val\_loss: 0.4440

Epoch 8/20

30/30 \_\_\_\_\_ 1s 37ms/step - accuracy: 0.9286 - loss: 0.3674 - val\_accuracy: 0.8732 - val\_loss: 0.4623

Epoch 9/20

30/30 \_\_\_\_\_ 1s 34ms/step - accuracy: 0.9290 - loss: 0.3720 - val\_accuracy: 0.8849 - val\_loss: 0.4455

Epoch 10/20

30/30 \_\_\_\_\_ 1s 36ms/step - accuracy: 0.9301 - loss: 0.3662 - val\_accuracy: 0.8835 - val\_loss: 0.4436

Epoch 11/20

30/30 \_\_\_\_\_ 2s 53ms/step - accuracy: 0.9248 - loss: 0.3696 - val\_accuracy: 0.8697 - val\_loss: 0.4647

Epoch 12/20

30/30 \_\_\_\_\_ 2s 36ms/step - accuracy: 0.9267 - loss: 0.3721 - val\_accuracy: 0.8829 - val\_loss: 0.4452

Epoch 13/20

30/30 \_\_\_\_\_ 1s 36ms/step - accuracy: 0.9334 - loss: 0.3637 - val\_accuracy: 0.8747 - val\_loss: 0.4587

Epoch 14/20

30/30 \_\_\_\_\_ 1s 36ms/step - accuracy: 0.9309 - loss: 0.3675 - val\_accuracy: 0.8837 - val\_loss: 0.4458

Epoch 15/20


30/30 \_\_\_\_\_ 1s 36ms/step - accuracy: 0.9337 - loss:

```
0.3599 - val_accuracy: 0.8799 - val_loss: 0.4553
Epoch 16/20
30/30 ━━━━━━━━━━━━━━━━━ 1s 37ms/step - accuracy: 0.9278 - loss:
0.3661 - val_accuracy: 0.8649 - val_loss: 0.4767
Epoch 17/20
30/30 ━━━━━━━━━━━━━━━━━ 1s 37ms/step - accuracy: 0.9349 - loss:
0.3637 - val_accuracy: 0.8813 - val_loss: 0.4474
Epoch 18/20
30/30 ━━━━━━━━━━━━━━━━━ 1s 37ms/step - accuracy: 0.9288 - loss:
0.3700 - val_accuracy: 0.8743 - val_loss: 0.4690
Epoch 19/20
30/30 ━━━━━━━━━━━━━━━━━ 1s 37ms/step - accuracy: 0.9259 - loss:
0.3697 - val_accuracy: 0.8679 - val_loss: 0.4798
Epoch 20/20
30/30 ━━━━━━━━━━━━━━━━━ 2s 47ms/step - accuracy: 0.9268 - loss:
0.3693 - val_accuracy: 0.8831 - val_loss: 0.4488
```

```
history_dict = history.history
history_dict.keys()
```

```
dict_keys(['accuracy', 'loss', 'val_accuracy', 'val_loss'])
```

```
import matplotlib.pyplot as plt
history_dict = history.history
loss_values = history_dict["loss"]
val_loss_values = history_dict["val_loss"]
epochs = range(1, len(loss_values) + 1)
plt.plot(epochs, loss_values, "bo", label="Training loss")
plt.plot(epochs, val_loss_values, "b", label="Validation loss")
plt.title("Training and validation loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.show()
```



```
plt.clf()
acc = history_dict["accuracy"]
val_acc = history_dict["val_accuracy"]
plt.plot(epochs, acc, "bo", label="Training acc")
plt.plot(epochs, val_acc, "b", label="Validation acc")
plt.title("Training and validation accuracy")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()
plt.show()
```

